



Dept. of Computer Science and Engineering  
University of Rajshahi  
www.ru.ac.bd

Dr. Shamim Ahmad

## Chapter 11: Storage and File Structure

- Brief overview of Physical Storage Media for Databases
  - ▶ To know its incidence on the design and usage of DBMSs
- Magnetic Disks
- RAID
- Storage Access and buffer management
- File Organization
  - Representation of records
  - Organization of Records in Files
  - Data-Dictionary Storage
- Storage and File Organization in Oracle 10g

José Aferes - Adaptado de Database System Concepts - 5<sup>th</sup> Edition 11.2

## Classification of Physical Storage Media

- In the end, a database must be physically stored in computer(s)
- Several aspects of storage media must be taken into account
  - Speed with which data can be accessed
  - Cost per unit of data
  - Reliability
    - ▶ data loss on power failure or system crash
    - ▶ physical failure of the storage device
  - Can differentiate storage into:
    - ▶ **volatile storage**: loses contents when power is switched off
    - ▶ **non-volatile storage**:
      - Contents persist even when power is switched off.
      - Includes secondary and tertiary storage, as well as better-backed up main-memory.

José Aferes - Adaptado de Database System Concepts - 5<sup>th</sup> Edition 11.3


## Physical Storage Media for DBMSs

- **Main memory**:
  - fast access (10s to 100s of nanoseconds; 1 nanosecond =  $10^{-9}$  seconds)
  - generally too small (or too expensive) to store the entire database
  - **Volatile** — contents of main memory are usually lost if a power failure or system crash occurs.
- **Magnetic-disk**
  - Primary medium for the long-term non-volatile online storage of data; typically stores entire database.
  - Data must be moved from disk to main memory for access, and written back for storage
    - ▶ Much slower access than main memory
  - **direct-access** — possible to read data on disk in any order
  - Reliable: usual mean time to failure is now of **about 3 to 5 years**
- **Optical and tape storage**
  - Mainly backups (offline storage)

José Aferes - Adaptado de Database System Concepts - 5<sup>th</sup> Edition 11.4

## Magnetic Hard Disk Mechanism

- It is worth taking a look at how magnetic disks work
  - After all they are the place where the databases are stored!



José Añeres - Adaptado de Database System Concepts - 5th Edition

115


## Magnetic Disks

- **Read-write head**
  - Positioned very close to the platter surface (almost touching it)
  - Reads or writes magnetically encoded information.
- Surface of platter divided into circular **tracks**
  - Over **50K-100K** tracks per platter on typical hard disks
- Each track is divided into **sectors**.
  - A sector is the smallest unit of data that can be read or written.
  - Sector size typically **512 bytes**
  - Typical sectors per track: **500** (on inner tracks) to **1000** (on outer tracks)
- To read/write a sector
  - disk arm swings to position head on right track
  - platter spins continually; data is read/written as sector passes under head
- Head-disk assemblies
  - multiple disk platters on a single **spindle** (**1 to 5 usually**)
  - one head per platter, mounted on a common arm.
- **Cylinder  $i$**  consists of  $i^{\text{th}}$  track of all the platters

José Añeres - Adaptado de Database System Concepts - 5th Edition

116

## Disk Subsystem



- Multiple disks connected to a computer system through a controller
  - **Controllers functionality** (checksum, bad sector remapping) often carried out by individual disks; reduces load on controller
- Disk interface standards families
  - **ATA (AT adaptor)** range of standards
  - **SATA (Serial ATA)**
  - **SCSI (Small Computer System Interconnect)** range of standards
  - Several variants of each standard (different speeds and capabilities)

José Añeres - Adaptado de Database System Concepts - 5th Edition

117

## Performance Measures of Disks

- **Access time** – the time it takes from when a read or write **request is issued** to when data transfer begins. Consists of:
  - **Seek time** – time it takes to **reposition the arm over the correct track**.
    - Average seek time is **1/2 the worst case seek time**.
    - **4 to 10 milliseconds** on typical disks
  - **Rotational latency** – time it takes for the sector to be accessed to appear under the head.
    - Average latency is **1/2 of the worst case latency**.
    - **4 to 11 milliseconds** on typical disks (**5400 to 15000 r.p.m.**)
- **Data-transfer rate** – the rate at which data can be retrieved from or stored to the disk
  - **25 to 100 MB per second** max rate, lower for inner tracks
  - Multiple disks may share a controller, so rate that controller can handle is also important
    - E.g. **ATA-5: 66 MB/sec, SATA: 150 MB/sec, Ultra 320 SCSI: 320 MB/s**
    - **Fiber Channel (FC2Gb): 256 MB/s**
- **Mean time to failure (MTTF)** – the average time the disk is expected to run continuously without any failure.
  - Nowadays, typically **3 to 5 years**

José Añeres - Adaptado de Database System Concepts - 5th Edition

118

## Optimization of Disk-Block Access

- **Block** – a contiguous sequence of sectors from a **single track**
  - data is transferred between disk and main memory in blocks
  - sizes range **from 512 bytes to several kilobytes**
    - **Smaller blocks**: more transfers from disk
    - **Larger blocks**: more space wasted due to partially filled blocks
    - Typical block sizes today range **from 4 to 16 kilobytes**
  - We'll see how this is important for database storage structure
- **Disk-arm-scheduling** algorithms order pending **accesses to tracks** so that disk arm movement is minimized
  - **elevator algorithm**: move disk arm in one direction (from outer to inner tracks or vice versa), **processing next request** in that direction, till **no more requests** in that direction, then reverse direction and repeat

## Optimization of Disk Block Access (Cont.)

- **File organization** – **optimize block access** time by organizing the blocks to correspond to how data will be accessed
  - E.g. Store related information on **the same or nearby cylinders**.
  - Files may get **fragmented** over time
    - E.g. if data is **inserted to/deleted** from the file
    - Or free blocks on disk are **scattered**, and newly created file has its blocks scattered over the disk
    - **Sequential access to a fragmented file** results in increased disk arm movement
  - Some systems have utilities to **defragment** the file system, in order to speed up file access

## RAID

- The choice of **disk structure** is very important in databases. Important factors, besides price, are:
  - **Capacity**
  - **Speed**
  - **Reliability**
- **RAID: Redundant Arrays of Independent Disks**
  - disk organization techniques that manage a large number of disks, providing a view of a single disk of
    - **high capacity** and **high speed** by using multiple disks in parallel, and
    - **high reliability** by storing data redundantly, so that data can be recovered even if a disk fails

## Improvement of Reliability via Redundancy

- **Redundancy** – store extra information that can be used to rebuild information lost in a disk failure
- E.g., **Mirroring (or shadowing)** – **RAID level 1**
  - Duplicate every disk. Logical disk consists of two physical disks.
  - Every write is carried out on both disks
    - Reads can take place from either disk
  - If one disk in a pair fails, data is still available in the other
    - Data loss would occur only if a disk fails, and its mirror disk also fails before the system is repaired
      - Probability of combined event is very small, (except for dependent failure modes such as fire or building collapse or electrical power surges)
- **Mean time to data loss** depends on mean time to failure, and **mean time to repair**


## Improvement in Performance via Parallelism

- Two main goals of parallelism in a disk system:
  1. Load balance multiple **small accesses** to increase throughput
  2. Parallelize large accesses to reduce response time.
- Improve transfer rate by **striping data** across multiple disks.
- **Bit-level striping** – split the bits of each byte across multiple disks
  - In an array of eight disks, write bit *i* of each byte to disk *i*.
  - Each access can read **data at eight times the rate of a single disk**.
  - But seek/access time worse than for a single disk
    - Bit level striping is not used much any more
- **Block-level striping** – with *n* disks, block *i* of a file goes to disk  $(i \bmod n) + 1$ 
  - Requests for different blocks can run in parallel if the blocks reside on different disks
  - A request for a **long sequence of blocks can utilize all disks in parallel**
  - Usually called **RAID level 0**

José Aferes - Adaptado de Database System Concepts - 5ª Edition 11.13

## RAID Levels 2 and 3

- **RAID Level 2: Memory-Style Error-Correcting-Codes (ECC)** with bit striping.
- **RAID Level 3: Bit-Interleaved Parity**
  - a single parity bit is enough for error correction, not just detection, since we know which disk has failed
    - When writing data, corresponding parity bits must also be computed and written to a parity bit disk
    - To recover data in a damaged disk, compute XOR of bits from other disks (including parity bit disk)
- RAID levels 4-6 elaborate on this idea of parity




José Aferes - Adaptado de Database System Concepts - 5ª Edition 11.14

## Choice of RAID Level

- Mirroring provides much better write performance than Parity
  - For Parity read operations are needed before a write, whereas Mirroring only requires 2 writes
  - Mirroring preferred for high update environments such as log disks
- Mirroring needs more disks for the same capacity
  - Higher cost per capacity unit
    - But price of storage capacity is already low and decreasing
- Nowadays, for database systems
  - (Distributed) Parity, with RAID Level 5, is preferred for applications with low update rate, and large amounts of data
  - Mirroring is preferred for all other applications
- The choice of the storage structure is an important first step when designing a real database!

José Aferes - Adaptado de Database System Concepts - 5ª Edition 11.15


## RAID – Redundant Array of Independent Disks



RAID Arrays

16

## RAID Array Components



Data Protection: RAID

17

## RAID Implementations

- ▶ **Hardware** (usually a specialized disk controller card)
  - Controls all drives attached to it
  - Array(s) appear to host operating system as a regular disk drive
  - Provided with administrative software
- ▶ **Software**
  - Runs as part of the operating system
  - Performance is dependent on CPU workload
  - Does not support all RAID levels

Data Protection: RAID

18

## RAID Levels


- ▶ 0 Striped array with no fault tolerance
- ▶ 1 Disk mirroring
- ▶ 3 Parallel access array with dedicated parity disk
- ▶ 4 Striped array with independent disks and a dedicated parity disk
- ▶ 5 Striped array with independent disks and distributed parity
- ▶ 6 Striped array with independent disks and dual distributed parity
- ▶ Nested RAID (i.e., 1 + 0, 0 + 1, etc.)

Data Protection: RAID

19



### RAID Redundancy: Parity



© 2008 EMC Corporation. All rights reserved.

RAID Arrays - 20

### Parity Calculation

$$5 + 3 + 4 + 2 = 14$$

The middle drive fails:

$$5 + 3 + ? + 2 = 14$$

$$? = 14 - 5 - 3 - 2$$

$$? = 4$$



© 2008 EMC Corporation. All rights reserved.

RAID Arrays - 21

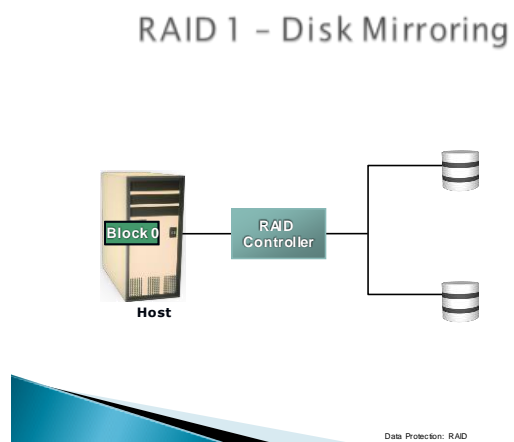
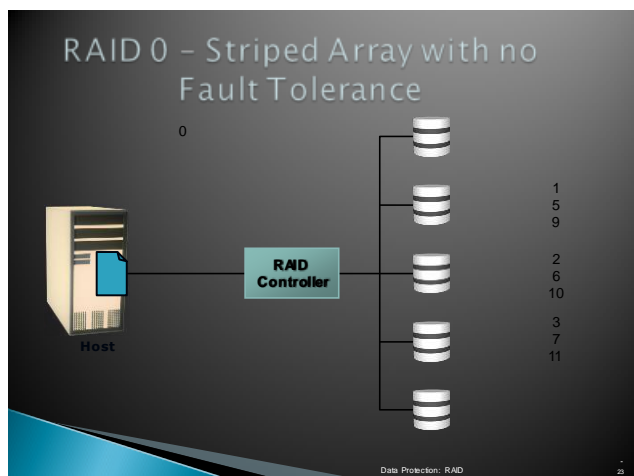


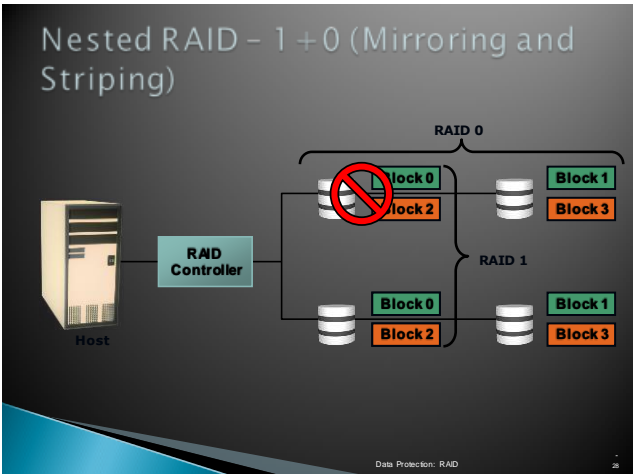
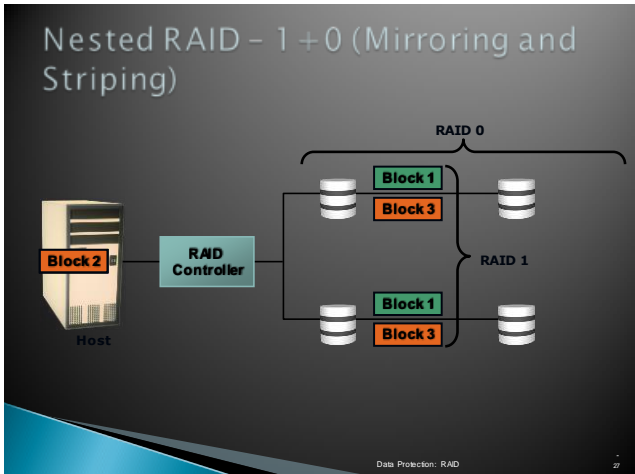
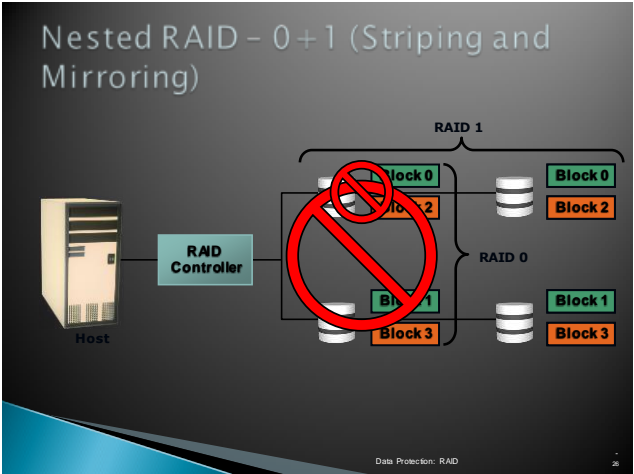
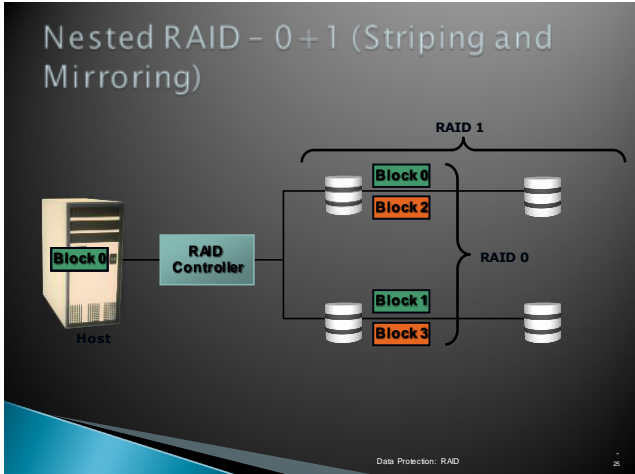
### Lecture 8, 9, 10

- Different RAID levels and their suitability for different application environments: RAID 0, RAID 1

© 2008 EMC Corporation. All rights reserved.

RAID Arrays - 22






## RAID 0+1 vs. RAID 1+0

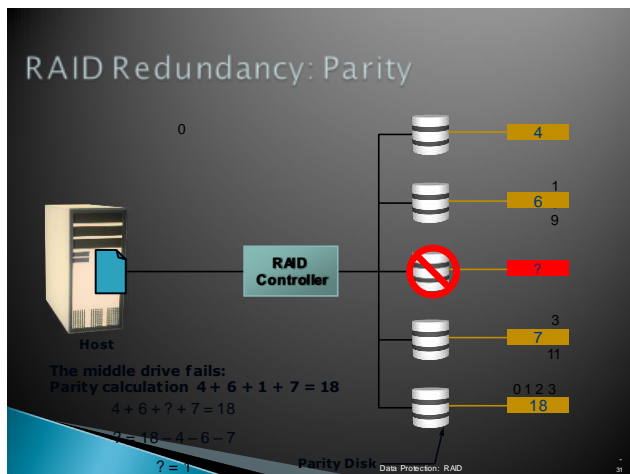
- Benefits are identical under normal operations
- Rebuild operations are very different
  - RAID 1+0 uses a mirrored pair – only 1 disk is rebuilt if a disk fails
  - RAID 0+1 if a single drive fails, the entire stripe is faulted
    - RAID 0+1 is a poorer solution and is less common

RAID Arrays


29



30




31



32



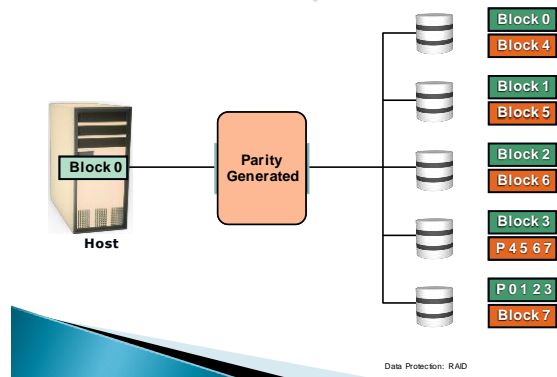
## RAID 4 – Striping with Dedicated Parity Disk



© 2008 EMC Corporation. All rights reserved.

RAID Arrays - 3D EMC

## RAID 5 – Independent Disks with Distributed Parity



Data Protection: RAID

34

## Buffer-Replacement Policies (Cont.)

- **Pinned block** – memory block that is not allowed to be written back to disk
- **Toss-immediate** strategy – frees the space occupied by a block as soon as the final tuple of that block has been processed
- **Most recently used (MRU) strategy** – system must pin the block currently being processed. After the final tuple of that block has been processed, the block is unpinned, and it becomes the most recently used block.
- Buffer managers also support **forced output** of blocks for the purpose of recovery
- To implement such specific policies a (good) database management system must usually implement its own buffer replacement policy (not relying in that of the operating system!)

José Aferes - Adaptado de Database System Concepts - 5<sup>th</sup> Edition 11.35

## File Organization

- The database is stored as a collection of *files*. Each file is a sequence of *records*. A record is a sequence of fields.
- First approach:
  - assume record size is fixed
  - each file has records of one particular type only
  - different files are used for different relations
 This case is the easiest to implement.

José Aferes - Adaptado de Database System Concepts - 5<sup>th</sup> Edition 11.36

## Fixed-Length Records

- Simple approach:
  - Store record  $i$  starting from byte  $n * (i - 1)$ , where  $n$  is the size of each record.
  - Record access is simple but records may cross blocks!
    - Modification: do not allow records to cross block boundaries

- Deletion of record  $i$ : alternatives:
  - move records  $i + 1, \dots, n$  to  $i, \dots, n - 1$
  - move record  $n$  to  $i$
  - do not move records, but link all free records on a *free list*

record 0	A-102	Perryridge	400
record 1	A-305	Round Hill	350
record 2	A-215	Mianus	700
record 3	A-101	Downtown	500
record 4	A-222	Redwood	700
record 5	A-201	Perryridge	900
record 6	A-217	Brighton	750
record 7	A-110	Downtown	600
record 8	A-218	Perryridge	700

## Free Lists


- Store the address of the first deleted record in the file header.
- Use this first record to store the address of the second deleted record, and so on
- Can think of these stored addresses as **pointers** since they "point" to the location of a record.
- More space efficient representation: reuse space for normal attributes of free records to store pointers. (No pointers stored in in-use records.)

header				
record 0	A-102	Perryridge	400	
record 1				
record 2	A-215	Mianus	700	
record 3	A-101	Downtown	500	
record 4				
record 5	A-201	Perryridge	900	
record 6				
record 7	A-110	Downtown	600	
record 8	A-218	Perryridge	700	

José Aferes - Adaptado de Database System Concepts - 5ª Edition 11.37

José Aferes - Adaptado de Database System Concepts - 5ª Edition 11.38

## Variable-Length Records: Slotted Page Structure



- Slotted page** are usually the size of a block
- Header contains:
  - number of record entries
  - end of free space in the block
  - location and size of each record
- Records can be moved around within a page to keep them contiguous with no empty space between them; entry in the header must be updated.
- (Other) pointers should not point directly to record — instead they should point to the entry for the record in header.

José Aferes - Adaptado de Database System Concepts - 5ª Edition 11.39

## Variable-Length Records

- Variable-length records arise in database systems in several ways:
  - Storage of multiple record types in a file.
  - Record types that allow variable lengths for one or more fields.
- If slotted pages are the size of a block, the issue of records spanning over more than one block is eliminated
- This limits the size of records in a database, which is usually the (default) case
  - There are special types for big records, that are treated differently (remember the **clob**s and **blob**s in Oracle?)

José Aferes - Adaptado de Database System Concepts - 5ª Edition 11.40

## Organization of Records in Files

- **Heap** – a record can be placed anywhere in the file where there is space
- **Sequential** – store records in sequential order, based on the value of the search key of each record
- **Hashing** – a hash function computed on some attribute of each record; the result specifies in which block of the file the record should be placed
- Records of each relation may be stored in a separate file. In a **multitable clustering file organization** records of several different relations can be stored in the same file
  - Motivation: store related records on the same block to minimize I/O
- The choice of a proper organization of records in a file is important for the efficiency of real databases!

José Aferes - Adaptado de Database System Concepts - 5ª Edition 11.41

## Sequential File Organization

- Suitable for applications that require sequential processing of the entire file
- The records in the file are ordered by a **search-key**

A-217	Brighton	750	
A-101	Downtown	500	
A-110	Downtown	600	
A-215	Mianus	700	
A-102	Perryridge	400	
A-201	Perryridge	900	
A-218	Perryridge	700	
A-222	Redwood	700	
A-305	Round Hill	350	

José Aferes - Adaptado de Database System Concepts - 5ª Edition 11.42

## Sequential File Organization (Cont.)

- Deletion – use pointer chains
- Insertion – locate the position where the record is to be inserted
  - if there is free space insert there
  - if no free space, insert the record in an **overflow block**
  - In either case, pointer chain must be updated
- Need to reorganize the file from time to time to restore sequential order

A-217	Brighton	750	
A-101	Downtown	500	
A-110	Downtown	600	
A-215	Mianus	700	
A-102	Perryridge	400	
A-201	Perryridge	900	
A-218	Perryridge	700	
A-222	Redwood	700	
A-305	Round Hill	350	
A-888	North Town	800	

José Aferes - Adaptado de Database System Concepts - 5ª Edition 11.43

## Multitable Clustering File Organization

Store several relations in one file using a **multitable clustering** file organization.

customer_name	account_number
Hayes	A-102
Hayes	A-220
Hayes	A-503
Turner	A-305

customer_name	customer_street	customer_city
Hayes	Main	Brooklyn
Turner	Putnam	Stamford

José Aferes - Adaptado de Database System Concepts - 5ª Edition 11.44

## Multitable Clustering File Organization (cont.)

Multitable clustering organization of *customer* and *depositor*:

Hayes	Main	Brooklyn
Hayes	A-102	
Hayes	A-220	
Hayes	A-503	
Turner	Putnam	Stamford
Turner	A-305	

- good for queries involving *depositor* ~~X~~ *customer*, and for queries involving one single customer and his accounts
- bad for queries involving only *customer*
  - but one can add pointer chains to link records of a particular relation
- results in variable size records

José Aferes - Adaptado de Database System Concepts - 5<sup>th</sup> Edition 11.45

## File System

- In sequential file organization, each relation is stored in a file
  - One may rely in the file system of the underlying operating system
- Multitable clustering may have significant gains in efficiency
  - But this may not be compatible with the file system of the operating system
- Several large scale database management systems do not rely directly on the underlying operating system
  - The relations are all stored in a single (multitable) file
  - The database management system manages the file by itself
  - This requires the implementation of an own file system inside the DBMS

José Aferes - Adaptado de Database System Concepts - 5<sup>th</sup> Edition 11.46

## Data Dictionary Storage

**Data dictionary** (also called **system catalog**) stores **metadata**; that is, data about data, such as

- Information about relations
  - names of relations
  - names and types of attributes of each relation
  - names and definitions of views
  - integrity constraints
- User and accounting information, including passwords
- Statistical and descriptive data
  - number of tuples in each relation
- Physical file organization information
  - How relation is stored (sequential/hash/...)
  - Physical location of relation
- Information about indices (next lecture)

José Aferes - Adaptado de Database System Concepts - 5<sup>th</sup> Edition 11.47

## Data Dictionary Storage (Cont.)

- Catalog structure
  - Relational representation on disk
  - specialized data structures designed for efficient access, in memory
- A possible catalog representation:

```

Relation_metadata = (relation_name, number_of_attributes,
                    storage_organization, location)
Attribute_metadata = (attribute_name, relation_name, domain_type,
                    position, length)
User_metadata = (user_name, encrypted_password, group)
Index_metadata = (index_name, relation_name, index_type,
                index_attributes)
View_metadata = (view_name, definition)
    
```

José Aferes - Adaptado de Database System Concepts - 5<sup>th</sup> Edition 11.48

## File Organization in Oracle

- Oracle has its own buffer management, with complex policies
- Oracle doesn't rely on the underlying operating system's file system
- A database in Oracle consists of **tablespaces**:
  - System tablespace: contains catalog meta-data
  - User data tablespaces
- The space in a tablespace is divided into **segments**:
  - Data segment
  - Index segment
  - Temporary segment (for sort operations)
  - Rollback segment (for processing transactions)
- Segments are divided into **extents**, each extent being a set of contiguous **database blocks**.
  - A database block need not be the same size of an operating system block, but is always a multiple

## File Organization in Oracle (cont.)

- A standard table is organized in a heap (no sequence is imposed)
- Partitioning of tables is possible for optimization
  - Range partitioning (e.g. by dates)
  - Hash partitioning
  - Composite partitioning
- Table data in Oracle can also be (multitable) clustered
  - One may tune the clusters to significantly improve the efficiency of query to frequently used joins.
- Hash file organization (to be studied in the sequel) is also possible for fetching the appropriate cluster
- A database can be tuned by an appropriate choice for the organization of data:
  - Choosing partitions
  - Appropriate choice of clusters
  - Hash or sequential
- Tuning makes the difference in big (real) databases!