



Dept. of Computer Science and Engineering
University of Rajshahi
www.ru.ac.bd

Dr. Shamim Ahmad

Cartesian-Product Operation – Example

Relations r, s :

A	B
α	1
β	2

r

C	D	E
α	10	a
β	10	a
β	20	b
γ	10	b

s

$r \times s$:

A	B	C	D	E
α	1	α	10	a
α	1	β	10	a
α	1	β	20	b
α	1	γ	10	b
β	2	α	10	a
β	2	β	10	a
β	2	β	20	b
β	2	γ	10	b

Cartesian-Product Operation

Notation $r \times s$

Defined as:

$$r \times s = \{tq \mid t \in r \text{ and } q \in s\}$$

Assume that attributes of $r(R)$ and $s(S)$ are disjoint. (That is, $R \cap S = \emptyset$).

If attributes of $r(R)$ and $s(S)$ are not disjoint, then renaming must be used.

If $r1(R1)$ and $r2(R2)$

then $r1 \times r2$ is a relation whose schema is the concatenation of $R1$ and $R2$.

Relation R contains all tuples t for which there is

a tuple $t1$ in $r1$ and

a tuple $t2$ in $r2$, for which

$t[R1] = t1[R1]$ and $t[R2] = t2[R2]$.

Composition of Operations

Can build expressions using multiple operations

Example: $\sigma_{A=C}(r \times s)$

$r \times s$

A	B	C	D	E
α	1	α	10	a
α	1	β	10	a
α	1	β	20	b
α	1	γ	10	b
β	2	α	10	a
β	2	β	10	a
β	2	β	20	b
β	2	γ	10	b

$\sigma_{A=C}(r \times s)$

A	B	C	D	E
1	α	10	a	
2	β	10	a	
2	β	20	b	

Rename Operation

- Allows us to name, and therefore to refer to, the results of relational-algebra expressions.
- Allows us to refer to a relation by more than one name.
- Example:

$$\rho_X(E)$$

returns the expression E under the name X

- If a relational-algebra expression E has arity n , then

$$\rho_{X(A_1, A_2, \dots, A_n)}(E)$$

returns the result of expression E under the name X , and with the attributes renamed to A_1, A_2, \dots, A_n .

Banking Example

branch (*branch_name*, *branch_city*, *assets*)

customer (*customer_name*, *customer_street*, *customer_city*)

account (*account_number*, *branch_name*, *balance*)

loan (*loan_number*, *branch_name*, *amount*)

depositor (*customer_name*, *account_number*)

borrower (*customer_name*, *loan_number*)

Example Queries

- Find all loans of over \$1200

$$\sigma_{amount > 1200}(loan)$$

- Find the loan number for each loan of an amount greater than \$1200

$$\Pi_{loan_number}(\sigma_{amount > 1200}(loan))$$

- Find the names of all customers who have a loan, an account, or both, from the bank

$$\Pi_{customer_name}(borrower) \cup \Pi_{customer_name}(depositor)$$

Example Queries

- Find the names of all customers who have a loan at the Perryridge branch.

$$\Pi_{customer_name}(\sigma_{branch_name = \text{"Perryridge"}}(\sigma_{borrower.loan_number = loan.loan_number}(borrower \times loan)))$$

- Find the names of all customers who have a loan at the Perryridge branch but do not have an account at any branch of the bank.

$$\Pi_{customer_name}(\sigma_{branch_name = \text{"Perryridge"}}(\sigma_{borrower.loan_number = loan.loan_number}(borrower \times loan))) - \Pi_{customer_name}(depositor)$$

Example Queries

- Find the names of all customers who have a loan at the Perryridge branch.

- Query 1

$$\Pi_{\text{customer_name}} (\sigma_{\text{branch_name} = \text{"Perryridge"}} (\sigma_{\text{borrower.loan_number} = \text{loan.loan_number}} (\text{borrower} \times \text{loan})))$$

- Query 2

$$\Pi_{\text{customer_name}} (\sigma_{\text{loan.loan_number} = \text{borrower.loan_number}} (\sigma_{\text{branch_name} = \text{"Perryridge"}} (\text{loan}) \times \text{borrower}))$$

Example Queries

- Find the largest account balance

- Strategy:

- Find those balances that are *not* the largest
 - Rename *account* relation as *d* so that we can compare each account balance with all others
- Use set difference to find those account balances that were *not* found in the earlier step.

- The query is:

$$\Pi_{\text{balance}}(\text{account}) - \Pi_{\text{account.balance}} (\sigma_{\text{account.balance} < \text{d.balance}} (\text{account} \times \rho_d(\text{account})))$$

Formal Definition

- A basic expression in the relational algebra consists of either one of the following:

- A relation in the database
- A constant relation

- Let E_1 and E_2 be relational-algebra expressions; the following are all relational-algebra expressions:

- $E_1 \cup E_2$
- $E_1 - E_2$
- $E_1 \times E_2$
- $\sigma_P(E_1)$, P is a predicate on attributes in E_1
- $\Pi_S(E_1)$, S is a list consisting of some of the attributes in E_1
- $\rho_X(E_1)$, X is the new name for the result of E_1

Additional Operations

We define additional operations that do not add any power to the relational algebra, but that simplify common queries.

- Set intersection
- Natural join
- Division
- Assignment

Set-Intersection Operation

- Notation: $r \cap s$
- Defined as:
- $r \cap s = \{ t \mid t \in r \text{ and } t \in s \}$
- Assume:
 - r, s have the same arity
 - attributes of r and s are compatible
- Note: $r \cap s = r - (r - s)$

Set-Intersection Operation – Example

- Relation r, s :

A	B
α	1
α	2
β	1

r

A	B
α	2
β	3

s

- $r \cap s$

A	B
α	2

Natural-Join Operation

- Notation: $r \bowtie s$
- Let r and s be relations on schemas R and S respectively. Then, $r \bowtie s$ is a relation on schema $R \cup S$ obtained as follows:
 - Consider each pair of tuples t_r from r and t_s from s .
 - If t_r and t_s have the same value on each of the attributes in $R \cap S$, add a tuple t to the result, where
 - ▶ t has the same value as t_r on r
 - ▶ t has the same value as t_s on s
- Example:
 - $R = (A, B, C, D)$
 - $S = (E, B, D)$
 - Result schema = (A, B, C, D, E)
 - $r \bowtie s$ is defined as:

$$\Pi_{r.A, r.B, r.C, r.D, s.E} (\sigma_{r.B=s.B \wedge r.D=s.D} (r \times s))$$

Natural Join Operation – Example

- Relations r, s :

A	B	C	D
α	1	α	a
β	2	γ	a
γ	4	β	b
α	1	γ	a
δ	2	β	b

r

B	D	E
1	a	α
3	a	β
1	a	γ
2	b	δ
3	b	ϵ

s

- $r \bowtie s$

A	B	C	D	E
α	1	α	a	α
α	1	α	a	γ
α	1	γ	a	α
α	1	γ	a	γ
δ	2	β	b	δ

Division Operation

- Notation: $r \div s$
- Suited to queries that include the phrase "for all".
- Let r and s be relations on schemas R and S respectively where

- $R = (A_1, \dots, A_m, B_1, \dots, B_n)$
- $S = (B_1, \dots, B_n)$

The result of $r \div s$ is a relation on schema

$R - S = (A_1, \dots, A_m)$

$$r \div s = \{ t \mid t \in \Pi_{R-S}(r) \wedge \forall u \in s (tu \in r) \}$$

Where tu means the concatenation of tuples t and u to produce a single tuple

Division Operation – Example

- Relations r, s :

A	B
α	1
α	2
α	3
β	1
γ	1
δ	1
δ	3
δ	4
ϵ	6
ϵ	1
β	2

B
1
2

s

- $r \div s$:

A
α
β

Another Division Example

- Relations r, s :

A	B	C	D	E
α	a	α	a	1
α	a	γ	a	1
α	a	γ	b	1
β	a	γ	a	1
β	a	γ	b	3
γ	a	γ	a	1
γ	a	γ	b	1
γ	a	β	b	1

r

D	E
a	1
b	1

s

- $r \div s$:

A	B	C
α	a	γ
γ	a	γ

Division Operation (Cont.)

- Property
 - Let $q = r \div s$
 - Then q is the largest relation satisfying $q \times s \subseteq r$
- Definition in terms of the basic algebra operation

Let $r(R)$ and $s(S)$ be relations, and let $S \subseteq R$

$$r \div s = \Pi_{R-S}(r) - \Pi_{R-S}((\Pi_{R-S}(r) \times s) - \Pi_{R-S,S}(r))$$

To see why

- $\Pi_{R-S,S}(r)$ simply reorders attributes of r
- $\Pi_{R-S}((\Pi_{R-S}(r) \times s) - \Pi_{R-S,S}(r))$ gives those tuples t in $\Pi_{R-S}(r)$ such that for some tuple $u \in s$, $tu \notin r$.



X1	X2	Y1	Y2
A	1	P	3
B	1	Q	2
A	1	P	3
C	2	P	1
A	1	P	1
B	1	Q	2
A	1	P	2
C	2	P	2
A	1	Q	2

Y1	Y2
P	1
P	2
Q	2



X1	X2
A	1
B	1
A	1
C	2
A	1
B	1
A	1
C	2
A	1

X1	X2
A	1
B	1
C	2

Y1	Y2
P	1
P	2
Q	2



X1	X2	Y1	Y2
A	1	P	1
B	1	P	2
C	2	Q	2

X1	X2	Y1	Y2
A	1	P	1
A	1	P	2
A	1	Q	2
B	1	P	1
B	1	P	2
B	1	Q	2
C	2	P	1
C	2	P	2
C	2	Q	2



X1	X2	Y1	Y2
A	1	P	1
A	1	P	2
A	1	Q	2
B	1	P	1
B	1	P	2
B	1	Q	2
C	2	P	1
C	2	P	2
C	2	Q	2

X1	X2	Y1	Y2
A	1	P	3
B	1	Q	2
A	1	P	3
C	2	P	1
A	1	P	1
B	1	Q	2
A	1	P	2
C	2	P	2
A	1	Q	2

X1	X2
B	1
C	2

Assignment Operation

- The assignment operation (\leftarrow) provides a convenient way to express complex queries.
 - Write query as a sequential program consisting of
 - a series of assignments
 - followed by an expression whose value is displayed as a result of the query.
 - Assignment must always be made to a temporary relation variable.
- Example: Write $r \div s$ as

$$\begin{aligned} temp1 &\leftarrow \Pi_{R,S}(r) \\ temp2 &\leftarrow \Pi_{R,S}((temp1 \times s) - \Pi_{R,S,S}(r)) \\ result &= temp1 - temp2 \end{aligned}$$
 - The result to the right of the \leftarrow is assigned to the relation variable on the left of the \leftarrow .
 - May use variable in subsequent expressions.

Bank Example Queries

- Find the names of all customers who have a loan and an account at bank

$$\Pi_{customer_name}(borrower) \cap \Pi_{customer_name}(depositor)$$

- Find the name of all customers who have a loan at the bank and the loan amount

$$\Pi_{customer_name, loan_number, amount}(borrower \bowtie loan)$$

Bank Example Queries

- Find all customers who have an account from at least the "Downtown" and the Uptown" branches.

- Query 1

$$\begin{aligned} &\Pi_{customer_name}(\sigma_{branch_name = "Downtown"}(depositor \bowtie account)) \cap \\ &\Pi_{customer_name}(\sigma_{branch_name = "Uptown"}(depositor \bowtie account)) \end{aligned}$$

- Query 2

$$\begin{aligned} &\Pi_{customer_name, branch_name}(depositor \bowtie account) \\ &\div \Pi_{temp(branch_name)}(((\text{"Downtown"}), (\text{"Uptown"}))) \end{aligned}$$

Note that Query 2 uses a constant relation.

Bank Example Queries

- Find all customers who have an account at all branches located in Brooklyn city.

$$\begin{aligned} &\Pi_{customer_name, branch_name}(depositor \bowtie account) \\ &\div \Pi_{branch_name}(\sigma_{branch_city = "Brooklyn"}(branch)) \end{aligned}$$

Extended Relational-Algebra-Operations

- Generalized Projection
- Aggregate Functions
- Outer Join

Generalized Projection

- Extends the projection operation by allowing arithmetic functions to be used in the projection list.

$$\Pi_{F_1, F_2, \dots, F_n}(E)$$

- E is any relational-algebra expression
- Each of F_1, F_2, \dots, F_n are arithmetic expressions involving constants and attributes in the schema of E .
- Given relation *credit_info*(*customer_name*, *limit*, *credit_balance*), find how much more each person can spend:

$$\Pi_{\text{customer_name}, \text{limit} - \text{credit_balance}}(\text{credit_info})$$

Aggregate Functions and Operations

- **Aggregation function** takes a collection of values and returns a single value as a result.

avg: average value
min: minimum value
max: maximum value
sum: sum of values
count: number of values

- **Aggregate operation** in relational algebra

$$G_1, G_2, \dots, G_n \mathcal{G}_{F_1(A_1), F_2(A_2), \dots, F_n(A_n)}(E)$$

E is any relational-algebra expression

- G_1, G_2, \dots, G_n is a list of attributes on which to group (can be empty)
- Each F_i is an aggregate function
- Each A_i is an attribute name

Aggregate Operation – Example

- Relation r :

A	B	C
α	α	7
α	β	7
β	β	3
β	β	10

- $\mathcal{G}_{\text{sum}(C)}(r)$

sum(C)
27

Aggregate Operation – Example

- Relation *account* grouped by *branch-name*:

<i>branch_name</i>	<i>account_number</i>	<i>balance</i>
Perryridge	A-102	400
Perryridge	A-201	900
Brighton	A-217	750
Brighton	A-215	750
Redwood	A-222	700

branch_name \mathcal{G} *sum(balance)* (*account*)

<i>branch_name</i>	<i>sum(balance)</i>
Perryridge	1300
Brighton	1500
Redwood	700

Aggregate Functions (Cont.)

- Result of aggregation does not have a name
 - Can use rename operation to give it a name
 - For convenience, we permit renaming as part of aggregate operation

branch_name \mathcal{G} *sum(balance)* **as** *sum_balance* (*account*)

Outer Join

- An extension of the join operation that avoids loss of information.
- Computes the join and then adds tuples from one relation that does not match tuples in the other relation to the result of the join.
- Uses *null* values:
 - *null* signifies that the value is unknown or does not exist
 - All comparisons involving *null* are (roughly speaking) **false** by definition.
 - We shall study precise meaning of comparisons with nulls later

Outer Join – Example

- Relation *loan*

<i>loan_number</i>	<i>branch_name</i>	<i>amount</i>
L-170	Downtown	3000
L-230	Redwood	4000
L-260	Perryridge	1700

- Relation *borrower*

<i>customer_name</i>	<i>loan_number</i>
Jones	L-170
Smith	L-230
Hayes	L-155

Outer Join – Example

Join

loan ⋈ *borrower*

<i>loan_number</i>	<i>branch_name</i>	<i>amount</i>	<i>customer_name</i>
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith

Left Outer Join

loan ⋈_L *borrower*

<i>loan_number</i>	<i>branch_name</i>	<i>amount</i>	<i>customer_name</i>
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith
L-260	Perryridge	1700	<i>null</i>

Outer Join – Example

Right Outer Join

loan ⋈_R *borrower*

<i>loan_number</i>	<i>branch_name</i>	<i>amount</i>	<i>customer_name</i>
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith
L-155	<i>null</i>	<i>null</i>	Hayes

Full Outer Join

loan ⋈_F *borrower*

<i>loan_number</i>	<i>branch_name</i>	<i>amount</i>	<i>customer_name</i>
L-170	Downtown	3000	Jones
L-230	Redwood	4000	Smith
L-260	Perryridge	1700	<i>null</i>
L-155	<i>null</i>	<i>null</i>	Hayes

Null Values

- It is possible for tuples to have a null value, denoted by *null*, for some of their attributes
- null* signifies an unknown value or that a value does not exist.
- The result of any arithmetic expression involving *null* is *null*.
- Aggregate functions simply ignore null values (as in SQL)
- For duplicate elimination and grouping, null is treated like any other value, and two nulls are assumed to be the same (as in SQL)

Null Values

- Comparisons with null values return the special truth value: *unknown*
- Three-valued logic using the truth value *unknown*:
 - OR: (*unknown* or *true*) = *true*,
(*unknown* or *false*) = *unknown*,
(*unknown* or *unknown*) = *unknown*
 - AND: (*true* and *unknown*) = *unknown*,
(*false* and *unknown*) = *false*,
(*unknown* and *unknown*) = *unknown*
 - NOT: (*not unknown*) = *unknown*
 - In SQL "*P* is *unknown*" evaluates to *true* if predicate *P* evaluates to *unknown*
- Result of select predicate is treated as *false* if it evaluates to *unknown*

Modification of the Database

- The content of the database may be modified using the following operations:
 - Deletion
 - Insertion
 - Updating
- All these operations are expressed using the assignment operator.

Deletion

- A delete request is expressed similarly to a query, except instead of displaying tuples to the user, the selected tuples are removed from the database.
- Can delete only whole tuples; cannot delete values on only particular attributes
- A deletion is expressed in relational algebra by:
$$r \leftarrow r - E$$
where r is a relation and E is a relational algebra query.

Deletion Examples

- Delete all account records in the Perryridge branch.
$$account \leftarrow account - \sigma_{branch_name = "Perryridge"}(account)$$
- Delete all loan records with amount in the range of 0 to 50
$$loan \leftarrow loan - \sigma_{amount \geq 0 \text{ and } amount \leq 50}(loan)$$
- Delete all accounts at branches located in Needham.
$$\begin{aligned} r_1 &\leftarrow \sigma_{branch_city = "Needham"}(account \bowtie branch) \\ r_2 &\leftarrow \pi_{account_number, branch_name, balance}(r_1) \\ r_3 &\leftarrow \pi_{customer_name, account_number}(r_2 \bowtie depositor) \\ account &\leftarrow account - r_2 \\ depositor &\leftarrow depositor - r_3 \end{aligned}$$

Insertion

- To insert data into a relation, we either:
 - specify a tuple to be inserted
 - write a query whose result is a set of tuples to be inserted
- in relational algebra, an insertion is expressed by:
$$r \leftarrow r \cup E$$
where r is a relation and E is a relational algebra expression.
- The insertion of a single tuple is expressed by letting E be a constant relation containing one tuple.

Insertion Examples

- Insert information in the database specifying that Smith has \$1200 in account A-973 at the Perryridge branch.

$$\begin{aligned} \text{account} &\leftarrow \text{account} \cup \{(\text{"A-973", "Perryridge", 1200})\} \\ \text{depositor} &\leftarrow \text{depositor} \cup \{(\text{"Smith", "A-973"})\} \end{aligned}$$

- Provide as a gift for all loan customers in the Perryridge branch, a \$200 savings account. Let the loan number serve as the account number for the new savings account.

$$\begin{aligned} r_1 &\leftarrow (\sigma_{\text{branch_name} = \text{"Perryridge"}}(\text{borrower loan})) \\ \text{account} &\leftarrow \text{account} \cup \Pi_{\text{loan_number}, \text{branch_name}, 200}(r_1) \\ \text{depositor} &\leftarrow \text{depositor} \cup \Pi_{\text{customer_name}, \text{loan_number}}(r_1) \end{aligned}$$

Updating

- A mechanism to change a value in a tuple without changing all values in the tuple
- Use the generalized projection operator to do this task

$$r \leftarrow \Pi_{F_1, F_2, \dots, F_i}(r)$$

- Each F_i is either
 - the i^{th} attribute of r , if the i^{th} attribute is not updated, or,
 - if the attribute is to be updated F_i is an expression, involving only constants and the attributes of r , which gives the new value for the attribute

Update Examples

- Make interest payments by increasing all balances by 5 percent.

$$\text{account} \leftarrow \Pi_{\text{account_number}, \text{branch_name}, \text{balance} * 1.05}(\text{account})$$

- Pay all accounts with balances over \$10,000 6 percent interest and pay all others 5 percent

$$\begin{aligned} \text{account} &\leftarrow \Pi_{\text{account_number}, \text{branch_name}, \text{balance} * 1.06}(\sigma_{\text{BAL} > 10000}(\text{account})) \\ &\cup \Pi_{\text{account_number}, \text{branch_name}, \text{balance} * 1.05}(\sigma_{\text{BAL} \leq 10000}(\text{account})) \end{aligned}$$
