



Dept. of Computer Science and Engineering
University of Rajshahi
www.ru.ac.bd

Dr. Shamim Ahmad

Transaction

Chapters 15: Transaction Management

- Transaction (Chapter 15)
 - Transaction Concept
 - Transaction State
 - Concurrent Executions
 - Serializability
 - Recoverability
 - Testing for Serializability

Transaction Concept

- A **transaction** is a *unit* of program execution that accesses and possibly updates various data items.
- E.g. transaction to transfer €50 from account A to account B:
 1. **read_from_account(A)**
 2. $A := A - 50$
 3. **write_to_account(A)**
 4. **read_from_account(B)**
 5. $B := B + 50$
 6. **write_to_account(B)**
- Two main issues to deal with:
 - Failures of various kinds, such as **hardware failures and system crashes**
 - Concurrent execution of **multiple transactions**

Transaction ACID properties

- E.g. transaction to transfer €50 from account A to account B:
 1. `read_from_account(A)`
 2. $A := A - 50$
 3. `write_to_account(A)`
 4. `read_from_account(B)`
 5. $B := B + 50$
 6. `write_to_account(B)`
- **Atomicity requirement**
 - if the transaction fails **after step 3 and before step 6**, money will be "lost" leading to an inconsistent database state
 - Failure could be due to software or hardware
 - the system should ensure that updates of a partially executed transaction are not reflected in the database
 - **All or nothing**, regarding the execution of the transaction
- **Durability requirement** — once the user has been notified of **transaction has completion**, the **updates must persist** in the database even if there are **software or hardware failures**.

José Aferes - Adaptado de Database System Concepts - 5th Edition 15-17.5

Transaction ACID properties (Cont.)

- Transaction to transfer €50 from account A to account B:
 1. `read_from_account(A)`
 2. $A := A - 50$
 3. `write_to_account(A)`
 4. `read_from_account(B)`
 5. $B := B + 50$
 6. `write_to_account(B)`
- **Consistency requirement** in above example:
 - **the sum of A and B** is unchanged by the execution of the transaction
- In general, consistency requirements include
 - Explicitly specified integrity constraints such as primary keys and foreign keys
 - Implicit integrity constraints
 - e.g. **sum of balances of all accounts**, minus sum of loan amounts must equal value of cash-in-hand
- A transaction must see a consistent database and must leave a consistent database
- During transaction execution the database may be **temporarily inconsistent**.
 - Constraints to be **verified only** at the end of the transaction

José Aferes - Adaptado de Database System Concepts - 5th Edition 15-17.6

Transaction ACID properties (Cont.)

- **Isolation requirement** — if between **steps 3 and 6**, another transaction T2 is allowed to access the **partially updated database**, it will see an inconsistent database (the sum $A + B$ will be less than it should be).
- | T1 | T2 |
|--------------------------|---|
| 1. <code>read(A)</code> | |
| 2. $A := A - 50$ | |
| 3. <code>write(A)</code> | |
| | <code>read(A), read(B), print(A+B)</code> |
| 4. <code>read(B)</code> | |
| 5. $B := B + 50$ | |
| 6. <code>write(B)</code> | |
- Isolation can be ensured trivially by running transactions **serially**
 - that is, one after the other.
 - However, executing multiple transactions concurrently has significant benefits, as we will see later.

José Aferes - Adaptado de Database System Concepts - 5th Edition 15-17.7

ACID Properties - Summary

A **transaction** is a **unit of program execution** that accesses and possibly updates various data items. To preserve the **integrity of data** the database system must ensure:

- **Atomicity** Either **all operations** of the transaction are properly reflected in the database or **none are**.
- **Consistency** Execution of a (single) transaction preserves the **consistency of the database**.
- **Isolation** Although multiple transactions may **execute concurrently**, each transaction must be unaware of other concurrently executing transactions. **Intermediate transaction results** must be hidden from other concurrently executed transactions.
 - That is, for every pair of transactions T_i and T_j , it appears to T_i that either T_j **finished** execution before T_i **started**, or T_j **started** execution after T_i **finished**.
- **Durability**. After a transaction completes successfully, the changes it has made to the **database persist, even if there are system failures**.

José Aferes - Adaptado de Database System Concepts - 5th Edition 15-17.8

Non-ACID Transactions

- There are application domains where ACID properties are not necessarily desired or, most likely, not always possible.
- This is the case of so-called **long-duration transactions**
 - Suppose that a transaction takes a lot of time
 - In this case it is **unlikely** that isolation can/should be guaranteed
 - ▶ E.g. Consider a transaction of booking a hotel and a flight
- Without Isolation, **Atomicity may be compromised**
- **Consistency and Durability should be preserved**
- Usual solution for long-duration transaction is to define **compensation action – what to do if later the transaction fails**
- In (centralized) databases long-duration transactions are usually not considered.
- But these are more and more important, specially in the **context of the Web**.

José Aferes - Adaptado de Database System Concepts - 5th Edition 15-17.9

Transaction State

- **Active** – the initial state; the transaction stays in this state while it is executing
- **Partially committed** – after the **final statement** has been executed.
- **Failed** – after the discovery that **normal execution can no longer proceed**.
- **Aborted** – after the transaction has been **rolled back** and the database **restored to its state prior** to the start of the transaction. Two options after it has been aborted:
 - **restart** the transaction
 - ▶ can be done only if no internal logical error
 - **kill** the transaction
- **Committed** – after **successful completion**.
- To guarantee atomicity, **external observable action** should all be performed (in order) after the **transaction is committed**.


José Aferes - Adaptado de Database System Concepts - 5th Edition 15-17.10

Committed Transactions

- A committed transaction that has performed update transforms the database into a new **consistent state**, which **must persist** even if there is a **system failure**.
- It cannot undo its effects by aborting it
- The only way to undo the effects of a committed transaction is to execute a **compensating transaction**.

José Aferes - Adaptado de Database System Concepts - 5th Edition 15-17.11


Transaction State (Cont.)



José Aferes - Adaptado de Database System Concepts - 5th Edition 15-17.12

Implementation of Atomicity and Durability


- The **recovery-management** component of a database system implements the support for atomicity and durability.
- E.g. the **shadow-database** scheme:
 - all updates are made on a **shadow copy** of the database
 - **db_pointer** is made to point to the updated shadow copy after
 - the transaction reaches partial commit and
 - all updated pages have been flushed to disk



José Aferes - Adaptado de Database System Concepts - 5ª Edition 15-17.13

Implementation of Atomicity and Durability

- The **recovery-management** component of a database system implements the support for atomicity and durability.
- E.g. the **shadow-database** scheme:
 - all updates are made on a **shadow copy** of the database
 - **db_pointer** is made to point to the updated shadow copy after
 - the transaction reaches partial commit and
 - after the operating system has written all the pages to disk



José Aferes - Adaptado de Database System Concepts - 5ª Edition 15-17.14

Implementation of Atomicity and Durability (Cont.)

- **db_pointer** always points to the current consistent copy of the database.
 - In case transaction fails, **old consistent** copy pointed to by **db_pointer** can be used, and the shadow copy can be deleted.
- The shadow-database scheme:
 - Assumes that **only one transaction** is active at a time.
 - Assumes **disks do not fail**
 - Useful for text editors, but
 - extremely **inefficient for large databases(!)**
 - Variant called shadow paging reduces copying of data, but is still not practical for large databases
 - Does not handle **concurrent transactions**
- Other implementations of atomicity and durability are possible, e.g. by using logs.
 - Log-based recovery will be addressed later.

José Aferes - Adaptado de Database System Concepts - 5ª Edition 15-17.15

Schedules

- **Schedule** – a sequence of instructions that specify the **chronological order** in which instructions of **concurrent transactions are executed**
 - a schedule for a **set of transactions** must consist of all instructions of those transactions
 - **must preserve the order** in which the instructions appear in each individual transaction.
- A transaction that successfully completes its execution will have a **commit instructions** as the last statement
 - by default transaction assumed to execute commit instruction as its last step
- A transaction that **fails to successfully** complete its execution will have an **abort instruction** as the last statement
- The goal is to find schedules that preserve the consistency.

José Aferes - Adaptado de Database System Concepts - 5ª Edition 15-17.16

Example Schedule 1

- Let T_1 transfer €50 from A to B , and T_2 transfer 10% of the balance from A to B .
- A **serial** schedule in which T_1 is followed by T_2 :

T_1	T_2
<code>read(A)</code> <code>A := A - 50</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + 50</code> <code>write(B)</code>	<code>read(A)</code> <code>temp := A * 0.1</code> <code>A := A - temp</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + temp</code> <code>write(B)</code>

José Aferes - Adaptado de Database System Concepts - 5th Edition 15-17.17

Example Schedule 2

- A serial schedule where T_2 is followed by T_1

T_1	T_2
<code>read(A)</code> <code>A := A - 50</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + 50</code> <code>write(B)</code>	<code>read(A)</code> <code>temp := A * 0.1</code> <code>A := A - temp</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + temp</code> <code>write(B)</code>

José Aferes - Adaptado de Database System Concepts - 5th Edition 15-17.18

Example Schedule 3

- Let T_1 and T_2 be the transactions defined previously. The following schedule is not a serial schedule, but it is **equivalent** to Schedule 1.

T_1	T_2
<code>read(A)</code> <code>A := A - 50</code> <code>write(A)</code>	<code>read(A)</code> <code>temp := A * 0.1</code> <code>A := A - temp</code> <code>write(A)</code>
<code>read(B)</code> <code>B := B + 50</code> <code>write(B)</code>	<code>read(B)</code> <code>B := B + temp</code> <code>write(B)</code>

In Schedules 1, 2 and 3, the sum $A + B$ is preserved.

José Aferes - Adaptado de Database System Concepts - 5th Edition 15-17.19

T_1	T_2
<code>read(A)</code> <code>write(A)</code>	<code>read(A)</code> <code>write(A)</code>
<code>read(B)</code> <code>write(B)</code>	<code>read(B)</code> <code>write(B)</code>

Schedule 3—showing only the read and write instructions

José Aferes - Adaptado de Database System Concepts - 5th Edition 15-17.20

Example Schedule 4

- The following concurrent schedule **does not preserve** the value of $(A + B)$.

T_1	T_2
read(A) $A := A - 50$	read(A) $temp := A * 0.1$ $A := A - temp$ write(A) read(B)
write(A) read(B) $B := B + 50$ write(B)	$B := B + temp$ write(B)

José Aferes - Adaptado de Database System Concepts - 5th Edition 15-17.21

Serializability

- Goal** : Deal with concurrent schedules that are equivalent to some serial execution:
 - Basic Assumption** – Each transaction preserves **database consistency**.
 - Thus **serial execution of a set of transactions** preserves database consistency.
- A (possibly concurrent) schedule is serializable if it is equivalent to a serial schedule. Different forms of schedule equivalence give rise to the notions of:
 - conflict serializability**
 - view serializability**
- Simplified view of transactions**
 - We **ignore operations** other than **read** and **write** instructions
 - We assume that transactions **may perform arbitrary computations** on data in **local buffers** in **between reads and writes**.
 - Our simplified schedules consist of only **read** and **write** instructions.

José Aferes - Adaptado de Database System Concepts - 5th Edition 15-17.22

Conflicting Instructions

- Instructions I_i and I_j of transactions T_i and T_j respectively, **conflict** if and only if there exists some item Q accessed by both I_i and I_j and at least **one of these instructions wrote Q** .
 - $I_i = \text{read}(Q)$, $I_j = \text{read}(Q)$. I_i and I_j don't conflict.
 - $I_i = \text{read}(Q)$, $I_j = \text{write}(Q)$. They conflict.
 - $I_i = \text{write}(Q)$, $I_j = \text{read}(Q)$. They conflict.
 - $I_i = \text{write}(Q)$, $I_j = \text{write}(Q)$. They conflict.
- Intuitively, a conflict between I_i and I_j forces an order between them.
 - If I_i and I_j are **consecutive in a schedule** and they do not conflict, their results would remain the same even if they had been interchanged in the schedule.

José Aferes - Adaptado de Database System Concepts - 5th Edition 15-17.23

T_1	T_2
read(A) write(A)	read(A) write(A)
read(B) write(B)	read(B) write(B)

Schedule 3—showing only the read and write instructions

José Aferes - Adaptado de Database System Concepts - 5th Edition 15-17.24

T_1	T_2
read(A)	
write(A)	
	read(A)
read(B)	
	write(A)
write(B)	
	read(B)
	write(B)

Figure 15.8 Schedule 5—schedule 3 after swapping of a pair of instructions.

Conflict Serializability

- If a schedule S can be transformed into a schedule S' by a series of swaps of non-conflicting instructions, we say that S and S' are **conflict equivalent**.
- We say that a schedule S is **conflict serializable** if it is conflict equivalent to a serial schedule.
- Schedule 3 can be transformed into Schedule 6, a serial schedule where T_2 follows T_1 , by series of swaps of non-conflicting instructions. Therefore it is conflict serializable.

T_1	T_2
read(A)	
write(A)	
	read(A)
	write(A)
read(B)	
write(B)	
	read(B)
	write(B)

Schedule 3

T_1	T_2
read(A)	
write(A)	
read(B)	
write(B)	
	read(A)
	write(A)
	read(B)
	write(B)

Schedule 6

José Aferes - Adaptado de Database System Concepts - 5th Edition 15-17.26

Conflict Serializability (Cont.)


- Example of a schedule-7 that is not conflict serializable:

T_3	T_4
read(Q)	
	write(Q)
write(Q)	

- We are unable to swap instructions in the above schedule to obtain either the serial schedule $\langle T_3, T_4 \rangle$, or the serial schedule $\langle T_4, T_3 \rangle$.

Testing for Serializability (Precedence Graph)

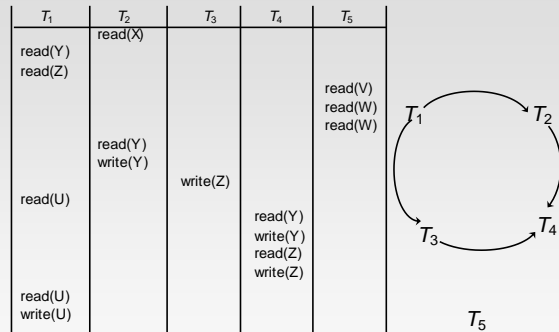
1. T_i executes write(Q) before T_j executes read(Q).
2. T_i executes read(Q) before T_j executes write(Q).
3. T_i executes write(Q) before T_j executes write(Q).



José Aferes - Adaptado de Database System Concepts - 5th Edition 15-17.27

José Aferes - Adaptado de Database System Concepts - 5th Edition 15-17.28

Example Schedule (Schedule A) + Precedence Graph



José Aferes - Adaptado de Database System Concepts - 5th Edition 15-17.29

View Serializability

- Sometimes it is possible to serialize schedules that are not conflict serializable
- View serializability provides a weaker and still consistency preserving notion of serialization
- Let S and S' be two schedules with the same set of transactions. S and S' are **view equivalent** if the following three conditions are met, for each data item Q ,
 - If in schedule S , transaction T_i reads the initial value of Q , then in schedule S' also transaction T_i must read the initial value of Q .
 - If in schedule S transaction T_i executes **read**(Q), and that value was produced by transaction T_j (if any), then in schedule S' also transaction T_i must read the value of Q that was produced by the same **write**(Q) operation of transaction T_j .
 - The transaction (if any) that performs the final **write**(Q) operation in schedule S must also perform the final **write**(Q) operation in schedule S' .

José Aferes - Adaptado de Database System Concepts - 5th Edition 15-17.30

T_1	T_5
read(A)	
$A := A - 50$	
write(A)	
	read(B)
	$B := B - 10$
	write(B)
read(B)	
$B := B + 50$	
write(B)	
	read(A)
	$A := A + 10$
	write(A)

Figure 15.11 Schedule 8.

View Serializability (Cont.)

- A schedule S is **view serializable** if it is view equivalent to a serial schedule.
- Every conflict serializable schedule is also view serializable.
- Below is a schedule which is view-serializable but *not* conflict serializable.

T_3	T_4	T_6
read(Q)		
write(Q)	write(Q)	
		write(Q)

- It is equivalent to either $\langle T_3, T_4, T_6 \rangle$ or $\langle T_4, T_3, T_6 \rangle$
- Every view serializable schedule that is not conflict serializable has **blind writes**.

José Aferes - Adaptado de Database System Concepts - 5th Edition 15-17.32